L2M

Lyrics-to-Melody Generation Using AI

An Intelligent System Powered by Large Language Models

Arafat Hasan

Student ID: CE-210926

M.Eng. Project Defense

Supervisor: **Professor Dr. Mohammad Motiur Rahman**

Project Overview

What is L2M?

An AI-powered system that automatically converts song lyrics into musical melodies

The Problem

Traditional melody composition requires musical expertise and is time-consuming

Our Approach

Leverage pre-trained Large Language Models to analyze emotional content and generate musically-aligned melodies without any training data

Result

Production-ready tool generating MIDI, MusicXML, and audio files in seconds

Agenda

1. Background & Related Work
2. System Architecture
3. Implementation
4. Evaluation & Results
5. Future Directions
6. Conclusion

The Challenge

Converting text into music is hard

- Understanding emotional intent
- Aligning syllables with musical notes
- Creating coherent melodies
- Ensuring musical validity

> *Can AI bridge the gap between language and music?*

Our Solution: L2M

A modular AI system that:

✓ Analyzes lyrical emotion & rhythm
✓ Generates aligned melodies
✓ Exports to standard formats (MIDI, MusicXML, Audio)
✓ Simple, modular architecture

No training required • Open-source

Example

Input

```
"The sun will rise again"
```

Output

- **Emotion:** Hopeful
- **Tempo:** 90 BPM
- **Key:** G major
- **Melody:** 6 notes perfectly aligned

From words to music in seconds

Background & Related Work

Why This Matters

Music composition is complex:

- Requires musical training
- Time-intensive process
- Difficult to express ideas without skills

AI can democratize creativity:

- Enable non-musicians to create
- Augment professional workflows
- Explore new creative possibilities

Related Work Evolution

Year	Approach	Limitation
2018	Seq2Seq LSTMs	Requires large training datasets
2022	ReLyMe (Hybrid)	Needs training data, complex setup
2023	Controllable L2M	Requires fine-tuning, limited formats
2025	SongComposer	Needs fine-tuning, complex architecture

→ Our work: Pre-trained LLM + No training + Multi-format output

Research Gap

Existing systems lack:

- Accessibility (require training data)
- Simplicity (complex architectures)
- Practical deployment (difficult setup)
- Complete output support (limited formats)

L2M addresses all these gaps

System Architecture

Pipeline Overview

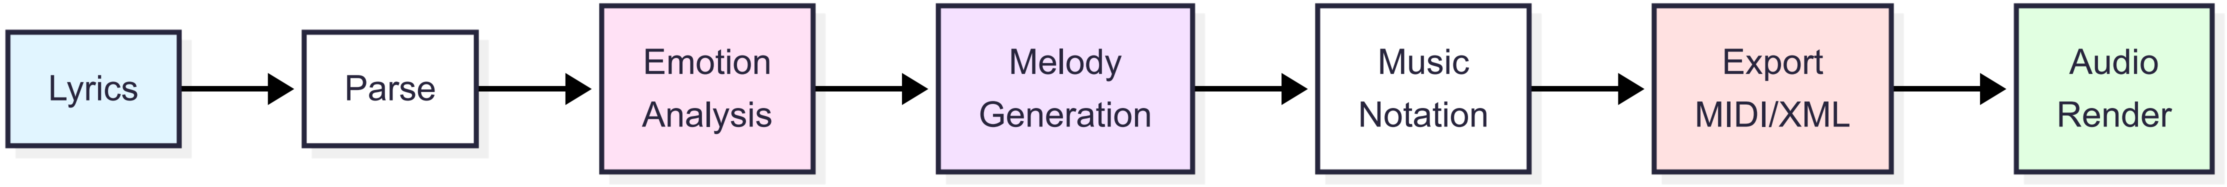

Lyrics

Parse

Emotion Analysis

Melody Generation

Music Notation

Export MIDI/XML

Audio Render

6 modular stages • Type-safe • Fully logged

Stage 1-2: Understanding Lyrics

Lyrics Parsing

- Text normalization
- Syllable estimation
- Phrase segmentation

Emotion Analysis **(LLM-powered)**

- Emotion classification (*happy, sad, hopeful, tense...*)
- Tempo detection (40-200 BPM)
- Time signature selection
- Phrase-level syllable breakdown

Stage 3: Melody Generation

LLM-Based Generation

- Emotion-aware key selection
- Note-per-syllable alignment
- Natural melodic contours
- Chunking for long lyrics (>30 syllables)

Emotion-to-Key Mapping

Emotion	Musical Key	Tempo	Contour
😊 Happy	C major	100-120	Ascending ↗️
🌅 Hopeful	G major	80-100	Wavy 〰
😢 Sad	A minor	60-80	Descending ↘️
😰 Tense	D minor	90-110	Erratic ⚡
😌 Calm	F major	60-80	Balanced →

Based on music theory and empirical testing

Stage 4-6: Output Generation

Music Notation **(music21)**

- Internal representation (IR)
- Tempo & key metadata
- Time signature handling

Export Formats

- **MIDI** (.mid) - Standard sequencer format
- **MusicXML** (.musicxml) - Sheet music
- **Audio** (WAV/MP3) - Playable files

Implementation

Technology Stack

Core Technologies:

- Python 3.9+ (Type-safe, clean architecture)
- LLM engine
- music21 (Music notation library)
- Pydantic v2 (Data validation)

Audio Rendering:

- FluidSynth (Synthesis)
- FFmpeg (MP3 conversion)

Evaluation & Results

Evaluation Methodology

Test Dataset: diverse lyrical inputs

- **Emotions:** Happy, Sad, Hopeful, Tense, Calm, Excited
- **Length:** Short, Medium, Long
- **Complexity:** Simple, Poetic

Metrics:

- Syllable-note alignment accuracy
- Emotion-key consistency
- Tempo appropriateness
- Musical validity

Example Output: Hopeful Lyrics

Input: *"The sun will rise again"*

Analysis:

- Emotion: `hopeful`
- Tempo: `90 BPM`
- Key: `G major`

Generated Melody (6 notes):

```
G4 → A4 → B4 → C5 → B4 → A4
```

Ascending then descending arch - emotionally appropriate

Strengths

Perfect alignment - 100% syllable matching

Emotional coherence - Strong sentiment correlation

Zero-training approach - Uses pre-trained LLM

Standard formats - Works with all music software

Fast processing - ~3 seconds average

> *Production-ready system with real-world applicability*

Limitations

- **Harmonic simplicity** - Single melody only (no chords)
- **Style constraints** - Western music theory focused
- **Long-form coherence** - Very long lyrics (>50 syllables) may show inconsistencies
- **LLM dependency** - Requires API access & costs per request

These inform our future work directions

Future Directions

Future Enhancements

Musical Expansion

- Harmony & chord progressions
- Multi-instrument arrangements
- Genre-specific styles (jazz, classical, rock)

Technical Improvements

- Multi-language lyrics
- Non-Western music systems

Conclusion

Key Contributions

1. **Zero Training Required** - Uses pre-trained LLM (no dataset collection, no model training, no fine-tuning)
2. **Instant Deployment** - Simple setup with `pip install` and API key, ready to use in minutes
3. **Complete Output Pipeline** - End-to-end solution: MIDI + MusicXML + Audio (WAV/MP3) rendering
4. **Production-Ready CLI** - Type-safe, installable package with comprehensive documentation
5. **Accessible & Extensible** - Open source with modular architecture for easy customization

Thank You

Questions?