



**L2M**

# **Lyrics-to-Melody Generation Using AI**

## **An Intelligent System Powered by Large Language Models**

**Arafat Hasan**

*M.Eng. Project Defense*

# **Agenda**

- 1. Background & Related Work**
- 2. System Architecture**
- 3. Implementation**
- 4. Evaluation & Results**
- 5. Future Directions**
- 6. Conclusion**

# The Challenge

**Converting text into music is hard**

- Understanding emotional intent
- Aligning syllables with musical notes
- Creating coherent melodies
- Ensuring musical validity

*| Can AI bridge the gap between language and music?*

# Our Solution: L2M

A modular AI system that:

- ✓ Analyzes lyrical emotion & rhythm
- ✓ Generates aligned melodies
- ✓ Exports to standard formats (MIDI, MusicXML, Audio)
- ✓ Simple, modular architecture

No training required • Open-source

# Example

## Input

"The sun will rise again"

## Output

- **Emotion:** Hopeful
- **Tempo:** 90 BPM
- **Key:** G major
- **Melody:** 6 notes perfectly aligned

## **Background & Related Work**

# Why This Matters

**Music composition is complex:**

- Requires musical training
- Time-intensive process
- Difficult to express ideas without skills

**AI can democratize creativity:**

- Enable non-musicians to create
- Augment professional workflows
- Explore new creative possibilities

## Related Work Evolution

Year	Approach	Limitation
2018	Seq2Seq LSTMs	Requires large training datasets
2022	ReLyMe (Hybrid)	Needs training data, complex setup
2023	Controllable L2M	Requires fine-tuning, limited formats
2025	SongComposer	Needs fine-tuning, complex architecture

→ Our work: Pre-trained LLM + No training + Multi-format output

# **Research Gap**

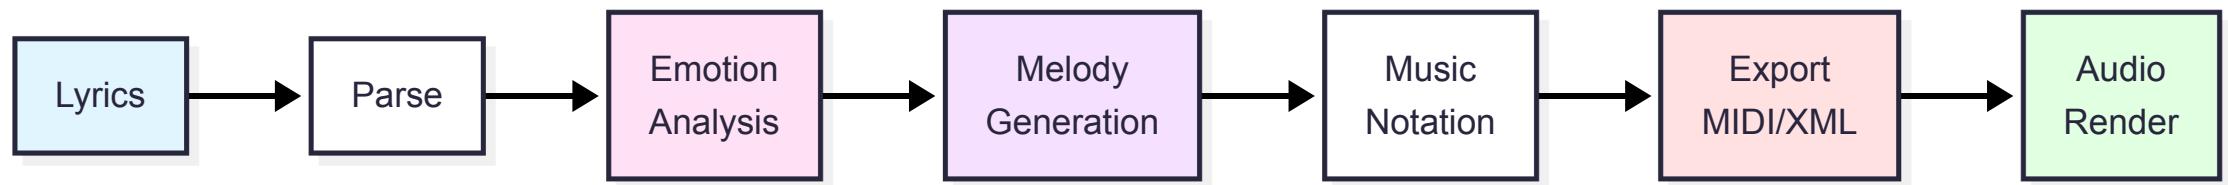
**Existing systems lack:**

- Accessibility (require training data)
- Simplicity (complex architectures)
- Practical deployment (difficult setup)
- Complete output support (limited formats)

**L2M addresses all these gaps**

# **System Architecture**

# Pipeline Overview



6 modular stages • Type-safe • Fully logged

# Stage 1-2: Understanding Lyrics

## Lyrics Parsing

- Text normalization
- Syllable estimation
- Phrase segmentation

## Emotion Analysis (LLM-powered)

- Emotion classification (*happy, sad, hopeful, tense...*)
- Tempo detection (40-200 BPM)
- Time signature selection

# Stage 3: Melody Generation

## LLM-Based Generation

- Emotion-aware key selection
- Note-per-syllable alignment
- Natural melodic contours
- Chunking for long lyrics (>30 syllables)

# Emotion-to-Key Mapping

Emotion	Musical Key	Tempo	Contour
 Happy	C major	100-120	Ascending 
 Hopeful	G major	80-100	Wavy ~
 Sad	A minor	60-80	Descending 
 Tense	D minor	90-110	Erratic 
 Calm	F major	60-80	Balanced →

*Based on music theory and empirical testing*

# Stage 4-6: Output Generation

## Music Notation (music21)

- Internal representation (IR)
- Tempo & key metadata
- Time signature handling

## Export Formats

- **MIDI** (.mid) - Standard sequencer format
- **MusicXML** (.musicxml) - Sheet music
- **Audio** (WAV/MP3) - Playable files

# Implementation

# Technology Stack

## Core Technologies:

- Python 3.9+ (Type-safe, clean architecture)
- LLM engine
- music21 (Music notation library)
- Pydantic v2 (Data validation)

## Audio Rendering:

- FluidSynth (Synthesis)
- FFmpeg (MP3 conversion)

# Key Components

## 1. LLMClient

Manages OpenAI API with retry logic

## 2. MelodyGenerator

Orchestrates melody creation with chunking

## 3. MIDIWriter

Converts IR to standard music formats

## 4. AudioRenderer

Synthesizes playable audio from MIDI

# Evaluation & Results

# Evaluation Methodology

**Test Dataset:** diverse lyrical inputs

- **Emotions:** Happy, Sad, Hopeful, Tense, Calm, Excited
- **Length:** Short, Medium, Long
- **Complexity:** Simple, Poetic

**Metrics:**

- Syllable-note alignment accuracy
- Emotion-key consistency
- Tempo appropriateness
- Musical validity

# Example Output: Hopeful Lyrics

Input: *"The sun will rise again"*

Analysis:

- Emotion: hopeful
- Tempo: 90 BPM
- Key: G major

Generated Melody (6 notes):

G4 → A4 → B4 → C5 → B4 → A4

*Ascending then descending arch - emotionally appropriate*

# Strengths

**Perfect alignment** - 100% syllable matching

**Emotional coherence** - Strong sentiment correlation

**Zero-training approach** - Uses pre-trained LLM

**Standard formats** - Works with all music software

**Fast processing** - ~3 seconds average

*Production-ready system with real-world applicability*

# Limitations

- **Harmonic simplicity** - Single melody only (no chords)
- **Style constraints** - Western music theory focused
- **Long-form coherence** - Very long lyrics (>50 syllables) may show inconsistencies
- **LLM dependency** - Requires API access & costs per request

*These inform our future work directions*

# Future Directions

# Future Enhancements

## Musical Expansion

- Harmony & chord progressions
- Multi-instrument arrangements
- Genre-specific styles (jazz, classical, rock)

## Technical Improvements

- Multi-language lyrics
- Non-Western music systems

# Conclusion

# Key Contributions

- 1. Zero Training Required** - Uses pre-trained LLM (no dataset collection, no model training, no fine-tuning)
- 2. Instant Deployment** - Simple setup with `pip install` and API key, ready to use in minutes
- 3. Complete Output Pipeline** - End-to-end solution: MIDI + MusicXML + Audio (WAV/MP3) rendering
- 4. Production-Ready CLI** - Type-safe, installable package with comprehensive documentation
- 5. Accessible & Extensible** - Open source with modular architecture for easy customization

**Thank You**

**Questions?**

